

Intervention-Conditioned Layer Fragility Across Language-Model Pretraining

Anonymous Authors

Abstract

Interventions on neural-network components are often interpreted as measuring properties such as redundancy or robustness, yet recent work shows that the conclusion can depend on the intervention protocol. We test whether such protocol dependence can be reduced to differences in output-level functional damage. Our assay measures the fraction of intact next-token top-1 predictions preserved after an interior-block intervention across independently pretrained small PYTHIA models. The early-to-late block-bypass response replicates across nine 160M runs, a held-out run, a 70M scale, and WikiText-2 and HellaSwag-derived evaluation streams. Within residual attenuation, matched comparisons show that KL/NLL predictive damage discriminates top-1 damage more strongly than raw activation displacement. However, prospective joint KL/NLL matching does not make block bypass and norm-controlled activation noise interchangeable: a corrected family residual of -0.0168 remains across five runs. A frozen block of simple output-logit and decision-boundary features accounts for only 20.2% of that residual. These results support a bounded methodological conclusion: in the tested setting, intervention responses are conditioned on how a model is perturbed and cannot be summarized by raw magnitude or matched predictive damage alone. We do not identify a causal mechanism, establish downstream capability loss, or claim generality beyond the tested small-model regime.

1 Introduction

Interventions are a central measurement tool for studying neural networks. Removing a layer, replacing its weights, patching an activation, or injecting noise can reveal whether a model’s output depends on the targeted component. The resulting change is often summarized as evidence about redundancy, necessity, robustness, or a causal role. This interpretation is strongest when different reasonable interventions agree. It is weaker when the response is a property of the measurement protocol as much as of the component being measured.

Recent work makes this concern concrete. Large language models can preserve a surprising fraction of their next-token predictions after layer deletion or swapping [7]. More directly, Garcia [5] shows that output-grounded layer equivalence varies across training and that replacement

and interchange need not agree. Related intervention audits likewise demonstrate that conclusions can depend on the comparator family and matching procedure [8]. We therefore do not ask whether protocol dependence exists or whether layer responses change over training. Those questions have substantial prior coverage.

We instead ask a narrower question: *can protocol-dependent responses be explained by differences in how much the interventions functionally damage the model?* This explanation is plausible. Two interventions with the same nominal strength can induce different changes in the output distribution; conversely, interventions that look different in activation or parameter space can be similar in function space. If protocol dependence disappears after matching forward KL and target-token NLL degradation, then intervention identity may add little beyond a severity budget. If it remains, intervention outcomes cannot be reduced to raw size or those output-level damage measures alone.

Our test builds on a deliberately narrow empirical base. We evaluate independently pretrained 70M- and 160M-parameter PYTHIA models [1]. For a fixed corpus and checkpoint, we intervene on one interior block and measure the fraction S of intact next-token top-1 predictions retained. The complement $D_S = 1 - S$ is top-1 damage, and $\Delta S = S_{\text{late}} - S_{\text{early}}$ summarizes the training endpoint change. This metric is not new: it belongs to the relative-accuracy family used for layer interventions [7] and is the complement of a divergent-token rate on identical positions [4]. Its role here is to provide a fixed outcome for controlled comparisons.

The block-bypass endpoint is reproducible: all nine 160M pretraining runs have negative ΔS on WikiText-2, including a prospectively held-out run, and the direction repeats at 70M and on a HellaSwag-derived evaluation stream. We use this result as an empirical substrate, not as the paper’s novelty claim. Within a continuous residual-attenuation family, magnitude-matched interventions retain a large functional-damage contrast, whereas the raw-magnitude contrast becomes small under joint KL/NLL matching. Functional damage is therefore the more informative descriptor under this assay, but it is not a randomized mediator.

The central experiment compares residual attenuation with norm-controlled additive activation noise. The noise family reproduces the qualitative training direction. We then target noise strengths prospectively, using only KL

and NLL, to match block interventions at the same run, checkpoint, and layer. After repairing and resealing a tie-breaking inconsistency, the canonical analysis retains a family residual of -0.0168 with a run-bootstrap 95% interval $[-0.0200, -0.0126]$: at matched predictive damage, activation noise produces less top-1 damage than block attenuation on average. A preregistered block of simple output-logit and decision-boundary features reduces this residual by 20.2%, but does not eliminate it.

Our contributions are exactly three:

1. We establish an independently replicated empirical base for a fixed block-bypass response across two small PYTHIA scales and two English evaluation streams, including held-out, confidence, and layer controls.
2. Through complementary magnitude-matched and damage-matched contrasts, we show that output-level predictive damage is more informative than raw activation displacement within the frozen residual-attenuation assay.
3. Through prospective joint KL/NLL targeting and a corrected deterministic top-1 analysis, we show that predictive-damage matching does not eliminate the difference between block attenuation and activation noise; simple output geometry explains only part of the remaining structure.

These are bounded empirical claims. We do not infer universal layer specialization, large-model scaling, downstream capability loss, KL/NLL causality, or an identified internal mechanism.

2 Related Work

Layer redundancy and protocol dependence. Layer deletion and swapping can preserve substantial next-token behavior in pretrained language models [7]. Block Influence and other depth-pruning criteria similarly show that some final-model layers can be removed with limited aggregate degradation [10]. Candidate A does not introduce layer deletion or top-1 agreement. Its closest neighbor is Garcia [5], who distinguishes replacement from interchange and shows that their output-grounded gap changes across Pythia checkpoints. That work directly occupies the broad claim that layer equivalence is protocol- and training-dependent. Our residual question is narrower: after matching two intervention families on output KL and NLL damage, does their top-1 response become interchangeable across independent pretraining runs?

Training-dependent perturbation fragility. Perturbation robustness is itself dynamic. Quantization degradation can vary substantially across pretraining and learning-rate phases [3]; activation-noise fragility can continue to resolve after probe accuracy saturates [11]. Layer-specific noise has also revealed heterogeneous learning dynamics in smaller neural networks [2]. These results preclude a generic novelty claim for training-dependent fragility. We use independently trained replicas to test sign stability

and focus the contribution on functional-damage matching between two fixed protocols.

Parameter, activation, and function space. Raw perturbation size does not generally determine functional change, because network composition, normalization, location, and direction transform local changes before they reach the output. Conversely, recent gradient-free language-model adaptation experiments identify norm as the dominant factor in their specific regime [6]. We treat this as a boundary condition: our data can reject a pure raw-displacement account within the frozen attenuation assay, not the relevance of norm in general. We measure function-space damage with forward KL from intact to intervened predictions and with intervened-minus-intact target-token NLL. These quantities share logits with top-1 damage, so alignment is descriptive unless a controlled contrast separates their predictions.

Mechanistic and causal interventions. Activation patching, ablation, and editing are widely used to localize computations, but their conclusions depend on corruption, metric, and intervention design [9, 13]. Self-repair can also compensate for ablated components in architecture- and prompt-specific ways [12]. SteerCheck provides the closest methodological collision: it prospectively matches off-target KL across activation-steering comparators and retains direction/family structure [8]. We do not claim matched-KL auditing as new. We apply a joint KL/NLL target to a different estimand—training-conditioned block substitutability—and make independent-run replication and assay repair part of the evidence contract.

Prediction-disagreement metrics. The divergent-token metrics of Deiseroth et al. [4] measure changes in argmax predictions under pruning and quantization. On the same teacher-forced positions, our agreement S is algebraically one minus such a disagreement rate. We use the statistic operationally and claim novelty only for the controlled result built around it.

3 Experimental Setup

Models and statistical units. We study independently pretrained PolyPythias replicas from the PYTHIA suite [1]. The principal empirical trajectory uses nine 160M runs at steps 14k, 36k, 72k, 107k, and 143k. A reused confirmation uses five 70M runs. Mechanism comparisons use six 160M runs (pilot labels 1, 4, and 9; confirmatory labels 6, 7, and 8) at steps 14k, 72k, and 143k. Cross-family targeting uses runs 2, 3, 5, 6, and 8 at those three checkpoints. The independent pretraining run is the scientific replicate; tokens, layers, matched pairs, and evaluation selections are not treated as independent replicates.

Evaluation streams. The primary stream consists of cached WikiText-2 training text. Three frozen selection seeds each produce six sequences with 256 predicted positions, yielding 4,608 next-token positions per checkpoint. Cross-corpus validation uses correct HellaSwag continuations transformed deterministically into a teacher-forced stream with the same tokenizer, selection algorithm, number of sequences, length, and positions. HellaSwag is an evaluation stream rather than a multiple-choice accuracy task in this assay.

Interventions. For block output $b(x)$ with residual input x , residual attenuation is

$$\tilde{b}_\alpha(x) = b(x) - \alpha [b(x) - x], \quad (1)$$

with $\alpha \in \{0.25, 0.50, 0.75, 1.00\}$. At $\alpha = 1$, the path returns x exactly and is the core block-bypass intervention. We intervene on blocks 1–10 of the 12-block 160M architecture. The second family adds a frozen random direction after the intact block and varies its scale β ; directions are averaged within a cell. Its scale is normalized by the local activation norm. The two families are intentionally structurally different.

Agreement and training response. For intact distribution p and intervened distribution q_I on fixed teacher-forced positions t , define

$$S(I, c) = \frac{1}{|T|} \sum_{t \in T} \mathbb{I}[\hat{y}_{c,t} = \hat{y}_{I,c,t}], \quad (2)$$

$$\hat{y}_{c,t} = \arg \max_y p_c(y \mid x_{<t}), \quad (3)$$

$$\hat{y}_{I,c,t} = \arg \max_y q_{I,c}(y \mid x_{<t}), \quad (4)$$

$$D_S(I, c) = 1 - S(I, c), \quad (5)$$

$$\Delta S(I) = S(I, c_{\text{late}}) - S(I, c_{\text{early}}). \quad (6)$$

The reported core statistic averages over frozen interior layers and evaluation selections. A negative ΔS means that the late checkpoint preserves fewer intact top-1 predictions under the same intervention.

Functional damage and raw magnitude. We measure

$$D_{\text{KL}}(I) = \mathbb{E}_t \text{KL}(p_t \| q_{I,t}), \quad (7)$$

$$D_{\text{NLL}}(I) = \text{NLL}(q_I) - \text{NLL}(p). \quad (8)$$

Raw magnitude is the local relative residual-change norm used in the frozen ARC-004 implementation; absolute activation RMS is a sensitivity measure. These measures are descriptors, not assumed causes.

Matching. Within residual attenuation, magnitude-matched pairs are formed within run and checkpoint with relative-magnitude ratio at most 1.10 and absolute KL separation at least 0.03. Damage-matched pairs require a KL

Table 1: Selected validated results. Intervals bootstrap independent-run summaries.

Result	Estimate	Replicates
WikiText-2 ΔS	−0.1035 [−0.1115, −0.0958]	9/9
HellaSwag ΔS	−0.0793 [−0.0885, −0.0677]	6/6
Magnitude-matched contrast	+0.0430 [0.0354, 0.0525]	6/6
Damage-matched contrast	+0.0025 [−0.0009, 0.0058]	4/6
Corrected family residual	−0.0168 [−0.0200, −0.0126]	5/5
Geometry-adjusted residual	−0.0134 [−0.0176, −0.0092]	5 runs

gap at most $\max(0.01, 0.10\bar{D}_{\text{KL}})$ and an NLL gap at most $\max(0.015, 0.10\bar{D}_{\text{NLL}})$, plus either a layer separation of at least three or a magnitude ratio of at least 1.25. Pairing is greedy, deterministic, and without reuse.

For the cross-family test, each block cell defines a joint KL/NLL target. A deterministic grid and bounded one-dimensional refinement select the noise strength using only strength, KL, NLL, and technical diagnostics. ‘Class A’ requires strict relative matching in both damage measures; ‘A+B’ is a declared sensitivity set. Top-1 outcomes remain hidden until the target manifest and analysis are sealed. The primary residual is

$$R_{\text{family}} = \text{mean}_{\text{run}} \text{median}_{i \in r} [D_S^{\text{noise}}(i) - D_S^{\text{block}}(i)]. \quad (9)$$

The frozen practical-equivalence band is ± 0.010742 .

Tie handling, confidence, and uncertainty. The canonical top-1 rule selects the lowest vocabulary index among exact maximum logits, implemented identically by ‘argmax’ for intact and both intervention families. Fixed intact-confidence bins are $[0, .05], [.05, .1], [.1, .2], [.2, .4], [.4, 1.01]$. Intervals are bootstrap intervals over independent-run summaries. All matching thresholds, exclusions, and primary contrasts were frozen before their corresponding confirmatory outcomes.

4 A Reproducible Intervention-Conditioned Response

We first establish that the outcome used for mechanism discrimination is not a single-checkpoint or single-run accident. This section is an empirical base for the functional-damage question; it is not presented as discovery of the broad training phenomenon.

Independent runs and held-out prediction. Across nine independent 160M pretraining runs, every late-minus-early block-bypass change is negative. Mean ΔS is −0.10353, with a run-bootstrap 95% interval $[-0.11149, -0.09585]$; the median is −0.10410. Run 9 was prospectively held out before its weights were evaluated. Its $\Delta S = -0.11895$ has the predicted sign and lies inside the frozen prediction interval $[-0.13463, -0.06366]$.

Intact NLL improves from an early-run mean of 3.878 to a late-run mean of 3.726, so the endpoint is not explained by a failure to learn next-token prediction.

The direction is broad across the tested locations: all 90 run–interior-layer endpoint changes are negative, and all ten layer medians are negative. The magnitude is not uniform; layer medians range from -0.0653 to -0.1838 . Accordingly, we claim broad sign support, not location invariance.

Confidence control. Top-1 agreement is mechanically related to intact confidence. We therefore repeat the endpoint comparison within five fixed intact-confidence bins. All 45 eligible run–bin changes are negative, with mean matched change -0.11388 . This control conditions on a coarse observable and does not fully match token difficulty.

Scale and corpus boundaries. Five independent 70M runs reproduce the direction (5/5 negative; median $\Delta S = -0.14334$, 95% interval $[-0.15385, -0.11030]$). This is replication at two small scales inside one architecture family, not a scaling law.

On the HellaSwag-derived stream, all six 160M runs are negative, with mean $\Delta S = -0.07926$ and 95% interval $[-0.08851, -0.06767]$. All 30 confidence- bin changes are also negative. The absolute effect is 0.766 times the WikiText-2 mean (95% interval $[0.645, 0.880]$): the direction repeats but the magnitude is corpus-dependent. HellaSwag is easier and more confident under the frozen stream, so the comparison does not imply corpus invariance.

Taken together, these checks justify using the endpoint response as a stable outcome for the controlled comparisons that follow. They do not establish why the response changes.

5 Functional Damage Explains More Than Raw Magnitude

We next discriminate between two simple descriptions of the response. A raw- magnitude account predicts that interventions moving the local activation by a similar amount should produce similar top-1 damage. A functional-damage account instead predicts that the change in output distribution is the more informative quantity. Because all variables are observational consequences of an intervention, we use matched contrasts rather than interpreting raw correlations.

5.1 Magnitude-matched comparison

The primary arm contains 185 disjoint pairs formed within run and checkpoint. Their median relative-magnitude ratio is 1.019, while the median absolute KL and NLL-damage separations are 0.087 and 0.098. In 94.6% of pairs, the higher-KL member also has higher NLL damage. The mean

run-median difference in top-1 damage is $+0.04297$ (95% interval $[0.03537, 0.05246]$), positive in all six runs and all three prospectively confirmatory runs. An absolute-RMS matching sensitivity is stronger rather than weaker.

Thus, approximately equal local displacement does not equalize the response when functional damage differs. The contrast holds within all five intact- confidence bins, although fixed bins are not exact propensity matches.

5.2 Damage-matched comparison

The reverse arm contains 162 disjoint pairs with median absolute KL and NLL gaps 0.0050 and 0.0070. Their median magnitude ratio is 1.129. The mean run-median top-1 contrast is $+0.00246$, with a 95% interval $[-0.00090, 0.00575]$. A stricter subset of 37 pairs with magnitude ratio at least 1.25 has contrast -0.00054 and interval $[-0.00456, 0.00371]$. Checkpoint-level signs in this stricter subset are heterogeneous.

The two contrasts have a clear ordering: in this attenuation family, functional-damage separation survives raw-magnitude matching, whereas raw-magnitude separation leaves little stable contrast after joint KL/NLL matching. This weakens a pure raw-displacement account. It does not show that raw magnitude is generally irrelevant: the observational regression retains a magnitude coefficient, location retains secondary structure, and the second intervention family behaves differently.

Finally, KL, NLL, and D_S are computed from the same intervened logits. The matched design shows discrimination among measured descriptors; it does not randomize KL/NLL or establish mediation.

6 Functional Damage Does Not Eliminate Intervention-Family Dependence

The within-family result motivates the paper’s central test. If output-level damage is an adequate description of an intervention, then two structurally different interventions matched on KL and NLL should have equivalent top-1 damage. We compare block attenuation with norm-controlled additive activation noise after the intact block.

Qualitative replication in a second family. The noise family reproduces the early-to-late direction: all five independent runs have negative ΔS , with mean -0.07246 and 95% interval $[-0.08438, -0.06304]$; all 25 eligible confidence-bin comparisons are negative. This rules out the narrow claim that the sign is unique to literal block bypass. It does not imply that the two families have the same quantitative response. An initial post-hoc cross-family comparison in ARC-005 had only 26/150 strict matches and failed its match-quality gate; its residual is not used as confirmatory evidence.




Figure 1: **Empirical base.** The endpoint direction repeats across independent runs, two small scales, and two evaluation streams. Panels reuse sealed ARC figures without recomputing outcomes. Differences in magnitude are not interpreted as invariance or a scaling law.




Figure 2: **Complementary matched contrasts within residual attenuation.** Functional-damage differences retain a stable top-1 contrast at matched magnitude, while the residual magnitude contrast is small after joint KL/NLL matching. Points and intervals summarize independent runs, not tokens.

Prospective joint-damage targeting. For each block target at the same run, checkpoint, and layer, a blinded one-dimensional search chooses the noise strength using only KL, NLL, strength, and technical diagnostics. It never computes or records D_S . Strict Class A support rises to 103/150 cells (68.7%); the A+B sensitivity set contains 134/150 (89.3%). Mean relative KL and NLL errors among Class A targets are 1.94% and 1.87%, respectively. Matching assignments and calipers are frozen before top-1 outcomes are revealed.

Assay invalidation and repair. The first reveal was invalid. Historical block outcomes had used ‘topk(2)[...,0]’, whereas new noise outcomes used ‘argmax’. Exact FP16 ties can resolve to different token indices; the maximum single-cell correction was comparable with the invalid aggregate. Analysis stopped before downstream geometry modeling. The original estimate is withdrawn.

ARC-006R changes only the identity convention: both families select the lowest vocabulary index among exact maximum logits. It does not change targets, matching classes, seeds, checkpoints, layers, equivalence bounds, exclusions, or the statistical protocol. The repair shifts

the invalid point estimate by +0.00240 (12.49% of its magnitude), demonstrating that the bug was scientifically material.


Corrected residual. Under the canonical rule,

$$R_{\text{family}} = -0.016833, \quad (10)$$


$$95\% \text{ CI} = [-0.020045, -0.012616]. \quad (11)$$

All five run medians are negative and every leave-one-run-out estimate remains below the negative equivalence boundary. Because the residual is $D_S^{\text{noise}} - D_S^{\text{block}}$, its negative sign means that block attenuation changes more top-1 predictions than matched activation noise on average, despite similar KL/NLL damage. The interval lies wholly outside the frozen ± 0.010742 equivalence band. The A+B sensitivity estimate is -0.01816 and is directionally consistent.

The aggregate is not cellwise universal: 24/103 corrected cells reverse sign, and layer means range from negative to positive. The result is therefore a run-level family component with substantial heterogeneity, not a deterministic law for every layer and checkpoint.



(a) Prospective matching support (outcome-independent)



(b) Corrected family residual under the canonical top-1 rule

Figure 3: **Predictive-damage matching does not eliminate family dependence.** Panel (a) is a matching diagnostic and remains valid because the assay repair did not change assignments. Panel (b) uses only corrected ARC-006R outcomes; the withdrawn mixed-rule estimate is not shown as evidence.

The safe conclusion is limited but consequential: in this assay, KL and NLL alignment is not sufficient to make the two interventions behaviorally exchangeable. The result does not identify which unmeasured property carries the difference.

7 Simple Output Geometry Is Not Sufficient

The family residual may still be a simple artifact of the top-1 decision rule. KL and NLL summarize the full predictive distribution, whereas D_S depends on whether a perturbation crosses the local argmax boundary. We therefore test a frozen sequence of non-tautological output-geometry models on the corrected 103-cell support set.

The baseline model conditions on the functional-damage summary and fixed run/checkpoint/layer effects. Adding the intact top-1 margin alone does not change the family coefficient. Adding anchored boundary displacement also fails to shrink it: the estimated magnitude grows by 1.49%. A full block that adds logit-change norm and absolute alignment reduces the family coefficient by 20.18%, from approximately -0.0168 to -0.01344 (95% interval $[-0.01764, -0.00923]$). The point estimate remains outside the frozen equivalence region, although its interval overlaps the boundary slightly.

The partial shrinkage is accompanied by counterevidence to a simple boundary story. An outcome-blind geometry-matched subset covers 31/103 cells and retains a residual of -0.02436 (95% interval $[-0.03006, -0.01819]$). The geometry features do not pass the frozen sign-reversal prediction gate (leave-one-run-out AUC 0.664), and layer heterogeneity increases by 26.2% after adjustment rather than decreasing. The far-margin stratum is near zero, but the effect is largest at medium rather than nearest margins.

We conclude that the tested output geometry is incomplete. We do not conclude that a deeper internal direction causes the remainder.

8 Discussion and Limitations

8.1 Interventions as conditional measurements


Our results support an interpretation of intervention outcomes as conditional measurements. Raw activation displacement alone is insufficient within the attenuation family, and output-level functional damage is more informative. Yet matching two intervention families on KL and NLL leaves a stable aggregate difference. Intervention identity therefore carries information not captured by these severity measures. This does not make family identity an explanation; it identifies the boundary of two simple explanations.

This distinction matters when an intervention response is promoted to a claim about a layer. A block may appear substitutable under one procedure and less so under another without contradiction: the procedures trace different paths through activation or parameter space. The appropriate scientific object is then the tuple of model, data, intervention, magnitude convention, output budget, and response metric—not “layer redundancy” without qualifiers.


8.2 What remains unidentified

Block attenuation and activation noise induce nearly orthogonal hidden perturbation directions, but direction disagreement varies too little across the matched cells to explain the outcome. Adding hidden direction features to the frozen output-geometry baseline improves held-out residual RMSE by 9.63%, below the preregistered 10% gate, and improves reversal AUC by only 0.0026. These analyses are exploratory.

The planned counterfactual direction test never reached outcome reveal. Its first calibration produced only 15/103 jointly supported cells. A continuous- α repair increased strict support to 47/103, still below the frozen



(a) Frozen model sequence and residual shrinkage



(b) Layer-level residual before and after geometry adjustment

Figure 4: **Output geometry is associated with only part of the family residual.** The full frozen geometry block reduces the aggregate coefficient, but the adjusted residual and heterogeneous layer structure remain. These are associational controls, not a causal decomposition.




Figure 5: **Evidence hierarchy, not a causal graph.** Arrows indicate the sequence of tested explanations. They do not assert mediation or causality.

minimum, with one layer receiving no support. The top-1 outcome remained blinded. Thus the data establish neither a zero nor a nonzero causal effect of internal direction. We stop this branch rather than relaxing support criteria after seeing feasibility results.

8.3 Limitations

Model scope. The core scale evidence is restricted to 70M and 160M PYTHIA replicas. The cross-family residual is evaluated only at 160M. We do not infer behavior at billion-parameter scale or in other architectures.

Data scope. The endpoint direction is tested on WikiText-2 and a HellaSwag-derived English stream with one tokenizer. Their magnitudes differ by roughly 23%. This is evidence against single-corpus specificity, not domain, language, or dataset universality. Cross-family damage matching is currently established on WikiText-2 only.

Intervention scope. Only residual attenuation/block bypass and norm-controlled additive activation noise re-

ceive the prospective cross-family test. Replacement, interchange, attention-only, MLP-only, and editing interventions may yield different relationships. Our result reinforces rather than removes the need to name the protocol.

Construct validity. S is teacher-forced next-token top-1 agreement. It is discontinuous near ties and does not measure semantic equivalence, downstream capability, or pruning safety. KL, NLL, geometry, and D_S share output logits, so even matched associations may reflect common metric geometry. The deterministic tie repair removes one inconsistency but does not make top-1 agreement a complete robustness construct.

Heterogeneity. Although the aggregate family residual is seed-stable, 24/103 cells reverse sign and layer effects vary substantially. Low support at layer 7 and sign changes across strata limit layer-specific interpretations.

Mechanism. The paper falsifies simple raw-magnitude, family-invariant damage, and simple output-geometry accounts under a frozen design. It does not identify the internal mechanism. This is a deliberate boundary, not evidence for whichever mechanism remains unnamed.

Assay repair. An inconsistent FP16 top-1 tie rule invalidated the first prospective reveal. We stopped, documented the issue, froze a family-independent lowest-index rule, and recomputed the analysis without changing matching or exclusions. The corrected qualitative conclusion survives, but the 12.49% point-estimate change demonstrates why the invalid result cannot be treated as harmless.

8.4 Implication

The strongest interpretation is methodological: when using internal interventions to infer component properties, researchers should report the intervention family, raw-space

convention, functional-damage budget, response metric, and support diagnostics. Agreement under one budget is not equivalent to protocol invariance, and residual protocol dependence should remain part of the claim rather than be averaged away.

9 Conclusion

We tested whether protocol-dependent layer-intervention responses can be reduced to raw perturbation magnitude or output-level predictive damage. A block-bypass top-1 end-point repeats across independent small-PYTHIA pretraining runs and two evaluation streams. Within residual attenuation, functional KL/NLL damage is more informative than raw activation displacement. Across block attenuation and norm-controlled activation noise, however, prospective joint KL/NLL matching leaves a corrected, seed-stable family residual. Simple output geometry accounts for only part of it, and the internal cause remains unidentified. The evidence therefore supports a bounded conclusion: intervention results in the tested setting are protocol-conditioned measurements, not protocol-independent layer properties.

References

- [1] Stella Biderman et al. Pythia: A suite for analyzing large language models across training and scaling. In *Proceedings of the 40th International Conference on Machine Learning*, 2023. URL <https://proceedings.mlr.press/v202/biderman23a.html>.
- [2] Hendrik Borras, Bernhard Klein, and Holger Froning. Walking noise: On layer-specific robustness of neural architectures against noisy computations and associated characteristic learning dynamics. *arXiv preprint arXiv:2212.10430*, 2022. URL <https://arxiv.org/abs/2212.10430>.
- [3] Albert Catalan-Tatjer, Niccolo Ajroldi, and Jonas Geiping. Training dynamics impact post-training quantization robustness. In *International Conference on Learning Representations*, 2026. URL <https://arxiv.org/abs/2510.06213>.
- [4] Bjorn Deiseroth, Max Meuer, Nikolas Gritsch, Constantin Eichenberg, Patrick Schramowski, Matthias Assenmacher, and Kristian Kersting. Divergent token metrics: Measuring degradation to prune away LLM components—and optimize quantization. In *Proceedings of NAACL-HLT*, 2024. URL <https://aclanthology.org/2024.naacl-long.377/>.
- [5] Gabriel Garcia. No free swap: Protocol-dependent layer redundancy in transformers. *arXiv preprint arXiv:2605.16234*, 2026. URL <https://arxiv.org/abs/2605.16234>.
- [6] Taeyeon Kim, Ahhyun Kim, TaeHyeon Kim, and Unggi Lee. Not the dimension, the norm: What matters in gradient-free weight perturbation of language models. *arXiv preprint arXiv:2608.01624*, 2026. URL <https://arxiv.org/abs/2608.01624>.
- [7] Vedang Lad, Jin Hwa Lee, Wes Gurnee, and Max Tegmark. The remarkable robustness of LLMs: Stages of inference? In *Advances in Neural Information Processing Systems*, 2025. URL https://proceedings.neurips.cc/paper_files/paper/2025/hash/bcad07d4bfab51243efaa08b8ed475b3-Abstract-Conference.html.
- [8] Daming Luo, Christy Liang, and Junyu Xuan. SteerCheck: Attribution specificity and alignment leakage in activation-steering audits. *arXiv preprint arXiv:2608.24335*, 2026. URL <https://arxiv.org/abs/2608.24335>.
- [9] Aleksandar Makelov, George Lange, Atticus Geiger, and Neel Nanda. Is this the subspace you are looking for? an interpretability illusion for subspace activation patching. In *International Conference on Learning Representations*, 2024. URL <https://openreview.net/forum?id=Ebt7Qd7Jwy>.
- [10] Xin Men, Mingyu Xu, Qingyu Zhang, Qianhao Yuan, Bingning Wang, Hongyu Lin, Yaojie Lu, Xianpei Han, and Weipeng Chen. ShortGPT: Layers in large language models are more redundant than you expect. In *Findings of the Association for Computational Linguistics: ACL 2025*, 2025. URL <https://aclanthology.org/2025.findings-acl.1035/>.
- [11] Orion Reblitz-Richardson. When probing accuracy saturates, fragility resolves: A complementary metric for LLM pre-training analysis. *arXiv preprint arXiv:2606.11375*, 2026. URL <https://arxiv.org/abs/2606.11375>.
- [12] Cody Rushing and Neel Nanda. Explorations of self-repair in language models. In *Proceedings of the 41st International Conference on Machine Learning*, 2024. URL <https://proceedings.mlr.press/v235/rushing24a.html>.
- [13] Fred Zhang and Neel Nanda. Towards best practices of activation patching in language models: Metrics and methods. *arXiv preprint arXiv:2309.16042*, 2023. URL <https://arxiv.org/abs/2309.16042>.

A Full Run Tables

The machine-readable source of every number appearing in the main text is ‘paper/numerical_provenance.csv’. Full run-level outputs remain in the sealed ARC directories rather than being copied or rounded again. The principal 160M run table is ‘meta_arc05/ARC-20260826-5060-003/results/run_level_results.csv’; the cross-corpus table is ‘research/arc.20260828_5060.011/processed/core_run_results.csv’.

B Seed-Level Robustness

The primary WikiText-2 endpoint is negative in 9/9 160M runs. All leave-one-run-out medians remain negative. The held-out ninth run is negative in all three evaluation selections. At 70M, 5/5 independent runs are negative. On the HellaSwag-derived stream, 6/6 runs and 18/18 fixed evaluation resamples are negative. For the corrected family residual, all five run medians are negative and all leave-one-run-out estimates remain below the negative equivalence boundary.

C Matching Diagnostics

Within-family matching is performed without pair reuse and within run and checkpoint. Primary magnitude matching yields 185 pairs; damage matching yields 162 pairs. Cross-family targeting obtains 103/150 strict Class A cells and 134/150 A+B cells. The strict-set mean relative KL and NLL target errors are 1.94% and 1.87%. Sixteen targets remain Class C. The worst target has 16.45% relative KL and 23.42% relative NLL error and is excluded by the frozen class rule, not by its outcome.

D Confidence Control

Fixed intact-confidence bins are specified in Section 3. The WikiText-2 core has 45/45 negative eligible run-bin endpoint changes. The activation-noise confirmation has 25/25 negative run-bin changes. The HellaSwag-derived stream has 30/30 negative changes. Fixed bins reduce a known composition difference but do not exactly match token difficulty.

E Cross-Corpus Details

Both streams use 4,608 predicted positions per checkpoint and the same tokenizer and sequence-selection procedure. WikiText-2 contains encyclopedia prose; the HellaSwag-derived stream contains correct everyday-event continuations. Mean intact NLL is lower and intact confidence higher on the latter. Its mean absolute endpoint effect is 0.766 of the WikiText-2 reference, with a bootstrap interval [0.645, 0.880]. The functional-damage versus raw-magnitude ordering repeats, but the damage-matched raw-magnitude contrast is small and above zero on HellaSwag. We therefore say “replicates across the tested corpora,” not “corpus invariant.”

F Assay Invalidation and Deterministic Top-1 Repair

ARC-006 compared historical block outcomes extracted with ‘topk’ against new noise outcomes extracted with ‘argmax’. Exact FP16 ties can choose different indices. ARC-007 detected the discrepancy at its reproduction gate before any geometry model was fit. The original -0.019235 residual and its reversal counts are invalid.

ARC-006R freezes a family-independent rule: select the lowest token index among exact maximum logits. It reuses the same 150 targets, classes, seeds, checkpoints, layers, calipers, and equivalence band. The corrected primary residual is -0.016833 ; a globally consistent ‘topk’ sensitivity is -0.016797 . Their difference is below the frozen numerical-instability threshold. The repair, source hashes, and output manifest are retained in the ARC-006R directory.

G Layer Heterogeneity

The core endpoint sign is negative for every run-layer pair, but magnitude is location dependent. In the corrected cross-family set, weighted layer-mean standard deviation is 0.01746 and the layer range crosses zero. Layer 7 has only

three strict matches and is marked low support. Output-geometry adjustment does not reduce this heterogeneity.

H Sign-Reversal Analysis

Twenty-four of 103 strict matched cells have $D_S^{\text{noise}} - D_S^{\text{block}} > 0$, despite the negative aggregate. A frozen geometry classifier reaches leave-one-run-out AUC 0.664 and does not pass the informativeness gate. These reversals are retained and prevent a uniform cellwise family claim.

I Output-Geometry Details

The preregistered model sequence adds intact margin, anchored boundary displacement, then logit-change norm and absolute alignment. KL and NLL are nearly collinear on the matched set, so the main regression uses a frozen functional-damage summary rather than interpreting their separate coefficients. The full geometry block shrinks the family coefficient by 20.18%; the geometry-matched subset and reversal analysis provide negative evidence against a complete output-boundary account.

J Internal-Direction Non-Identifiability

The hidden perturbation vectors for the two families are almost orthogonal. Observational direction features miss the frozen held-out improvement gates. More importantly, the planned counterfactual test fails support before outcome reveal. The paper therefore makes no statement about the sign or magnitude of a directional causal effect.

K Continuous-Alpha Feasibility Failure

ARC-009 uses continuous alpha to improve joint support without exposing D_S . Strict PASS-A support rises from 15 to 47 cells and PASS-B reaches 29, but the frozen minimum is not met and layer 2 has 0/15 PASS-A cells. The causal branch terminates without a reveal; thresholds are not relaxed post hoc.

L Reproducibility and Provenance

Every ARC stores preregistration or protocol documents, raw/processed results, run manifests, commands, integrity hashes, and resource accounting. The Draft 0 provenance table records the exact artifact used for each number. GPU inference is not rerun for manuscript construction. The empirical chain follows the supersession rule ARC-006 $\rightarrow$ ARC-006R; later analyses reproduce the corrected baseline before proceeding.